

A photograph of a university campus featuring a row of modern, multi-story concrete buildings with balconies. In the foreground, there is a paved area with a circular fountain in the center. Lush green trees and bushes are scattered throughout the scene. The sky is filled with soft, white clouds. A semi-transparent white box is overlaid on the bottom left of the image, containing text.

CSC2B10/CSC02B2

Data Communications

Chapter 3.1



UNIVERSITY
OF
JOHANNESBURG

Chapter 3: Transport Layer

Our goals:

- understand principles behind transport layer services:
 - multiplexing/demultiplexing
 - reliable data transfer
 - flow control
 - congestion control
- learn about transport layer protocols in the Internet:
 - UDP: connectionless transport
 - TCP: connection-oriented transport
 - TCP congestion control

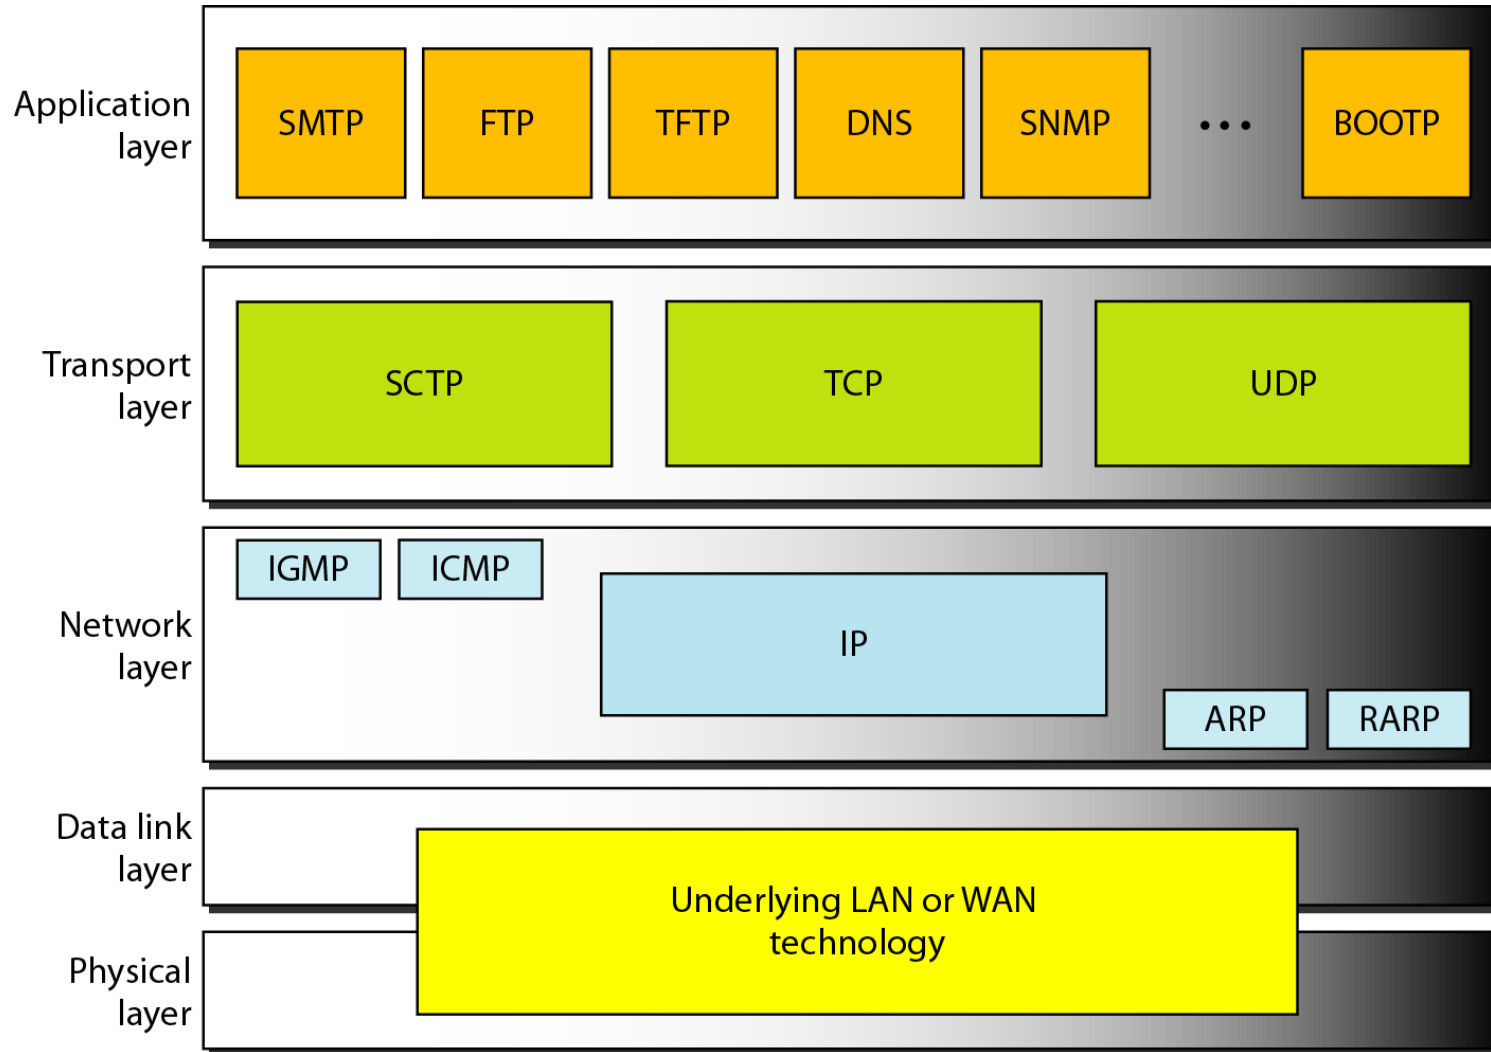


Transport layer: roadmap

- Transport-layer services
 - Multiplexing and demultiplexing
 - Connectionless transport: UDP
 - Principles of reliable data transfer
 - Connection-oriented transport: TCP
 - Principles of congestion control
 - TCP congestion control
 - Evolution of transport-layer functionality
-

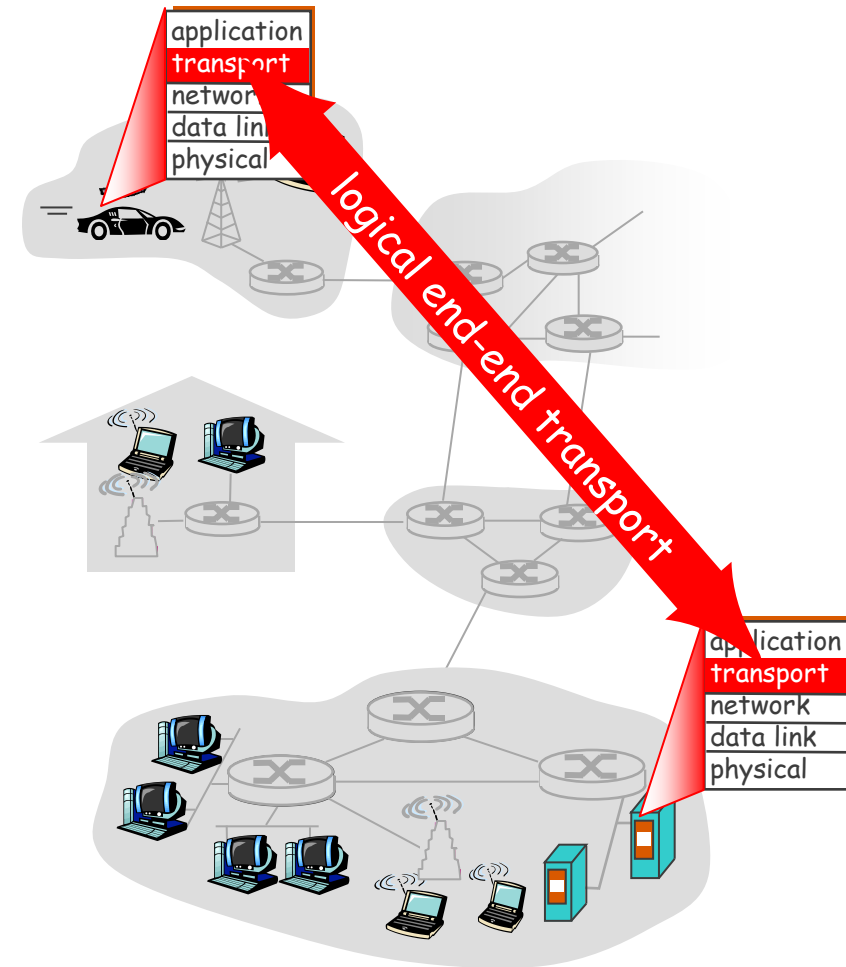


Position of UDP, TCP, and SCTP in TCP/IP suite



Transport services and protocols

- provide *logical communication* between app processes running on different hosts
- transport protocols run in end systems
 - send side: breaks app messages into **segments**, passes to network layer
 - rcv side: reassembles segments into messages, passes to app layer
- more than one transport protocol available to apps
 - Internet: **TCP and UDP**



Transport vs. network layer services and protocols



- **network layer:**
logical
communication
between hosts
- **transport layer:**
logical
communication
between processes
 - relies on, enhances,
network layer
services

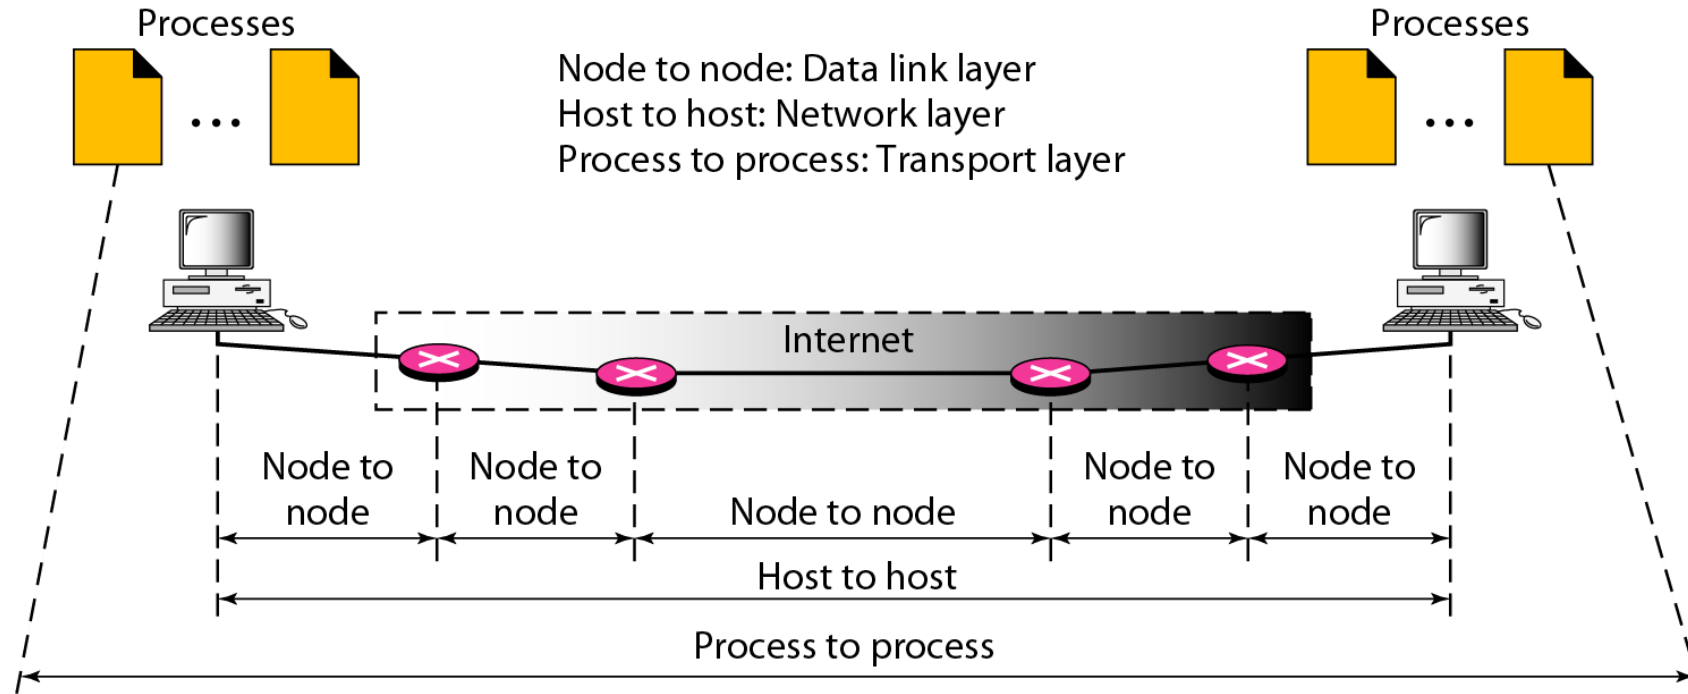
household analogy:

*12 kids in Ann's house sending
letters to 12 kids in Bill's
house:*

- hosts = houses
- processes = kids
- app messages = letters in envelopes



Process to process delivery



Transport vs. network layer

Addressing:

Data link – MAC address

Network – IP address

Transport – port number

16 bit integer – 0-65535

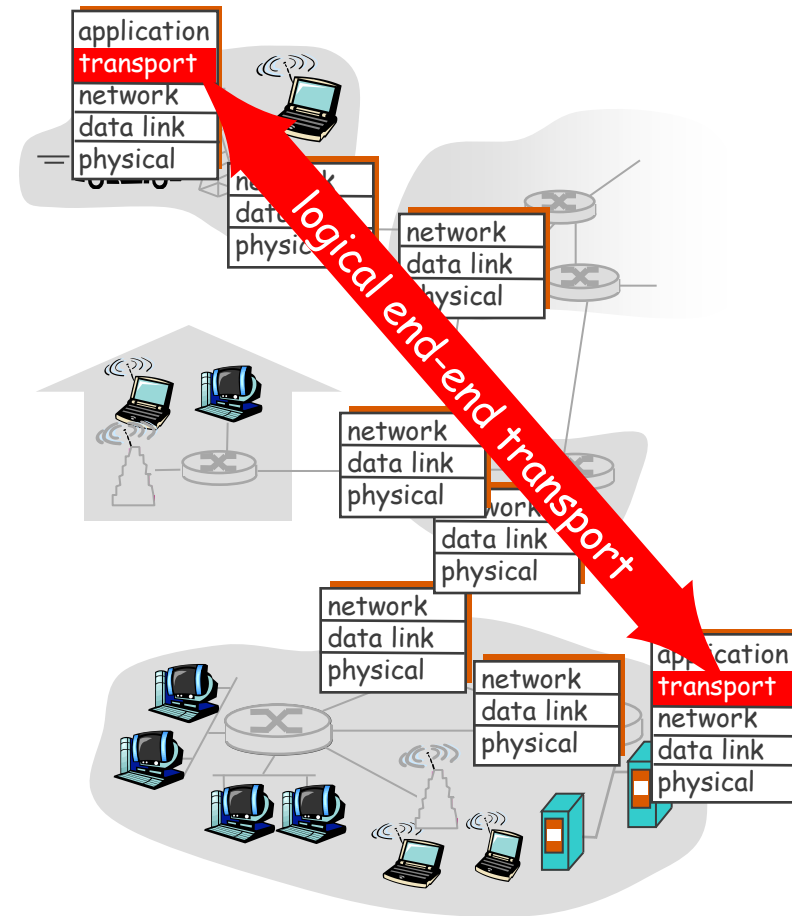
Client port – ephemeral port number

Server port – not random



Internet transport-layer protocols

- reliable, in-order delivery (**TCP**)
 - congestion control
 - flow control
 - connection setup
- unreliable, unordered delivery: **UDP**
 - no-frills extension of “best-effort” IP
- services not available:
 - delay guarantees
 - bandwidth guarantees



Chapter 3 outline

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

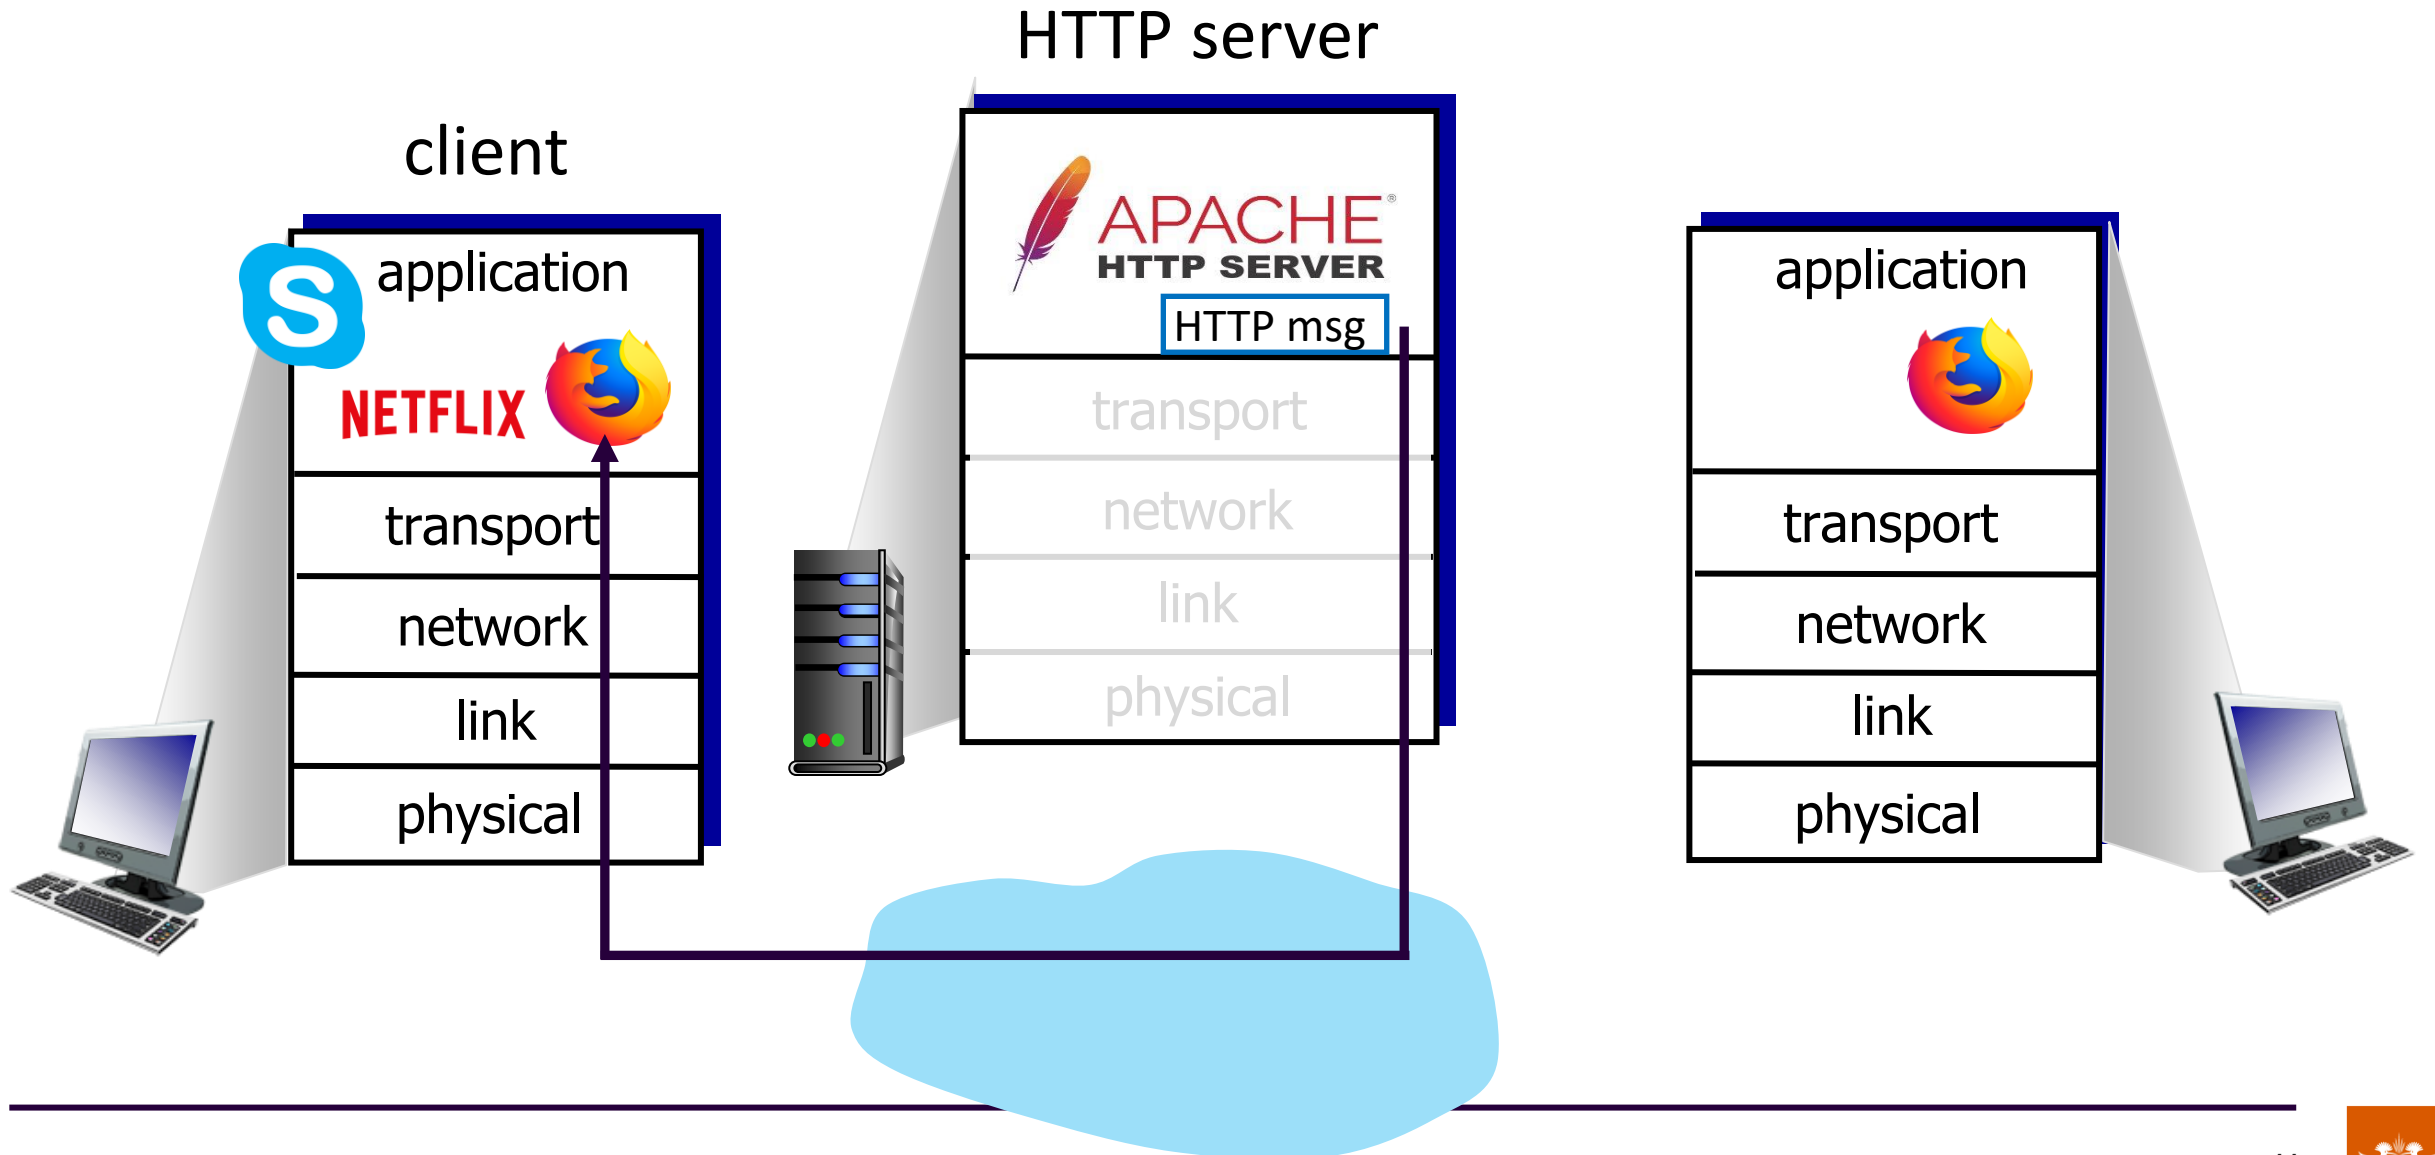
3.5 Connection-oriented transport:
TCP

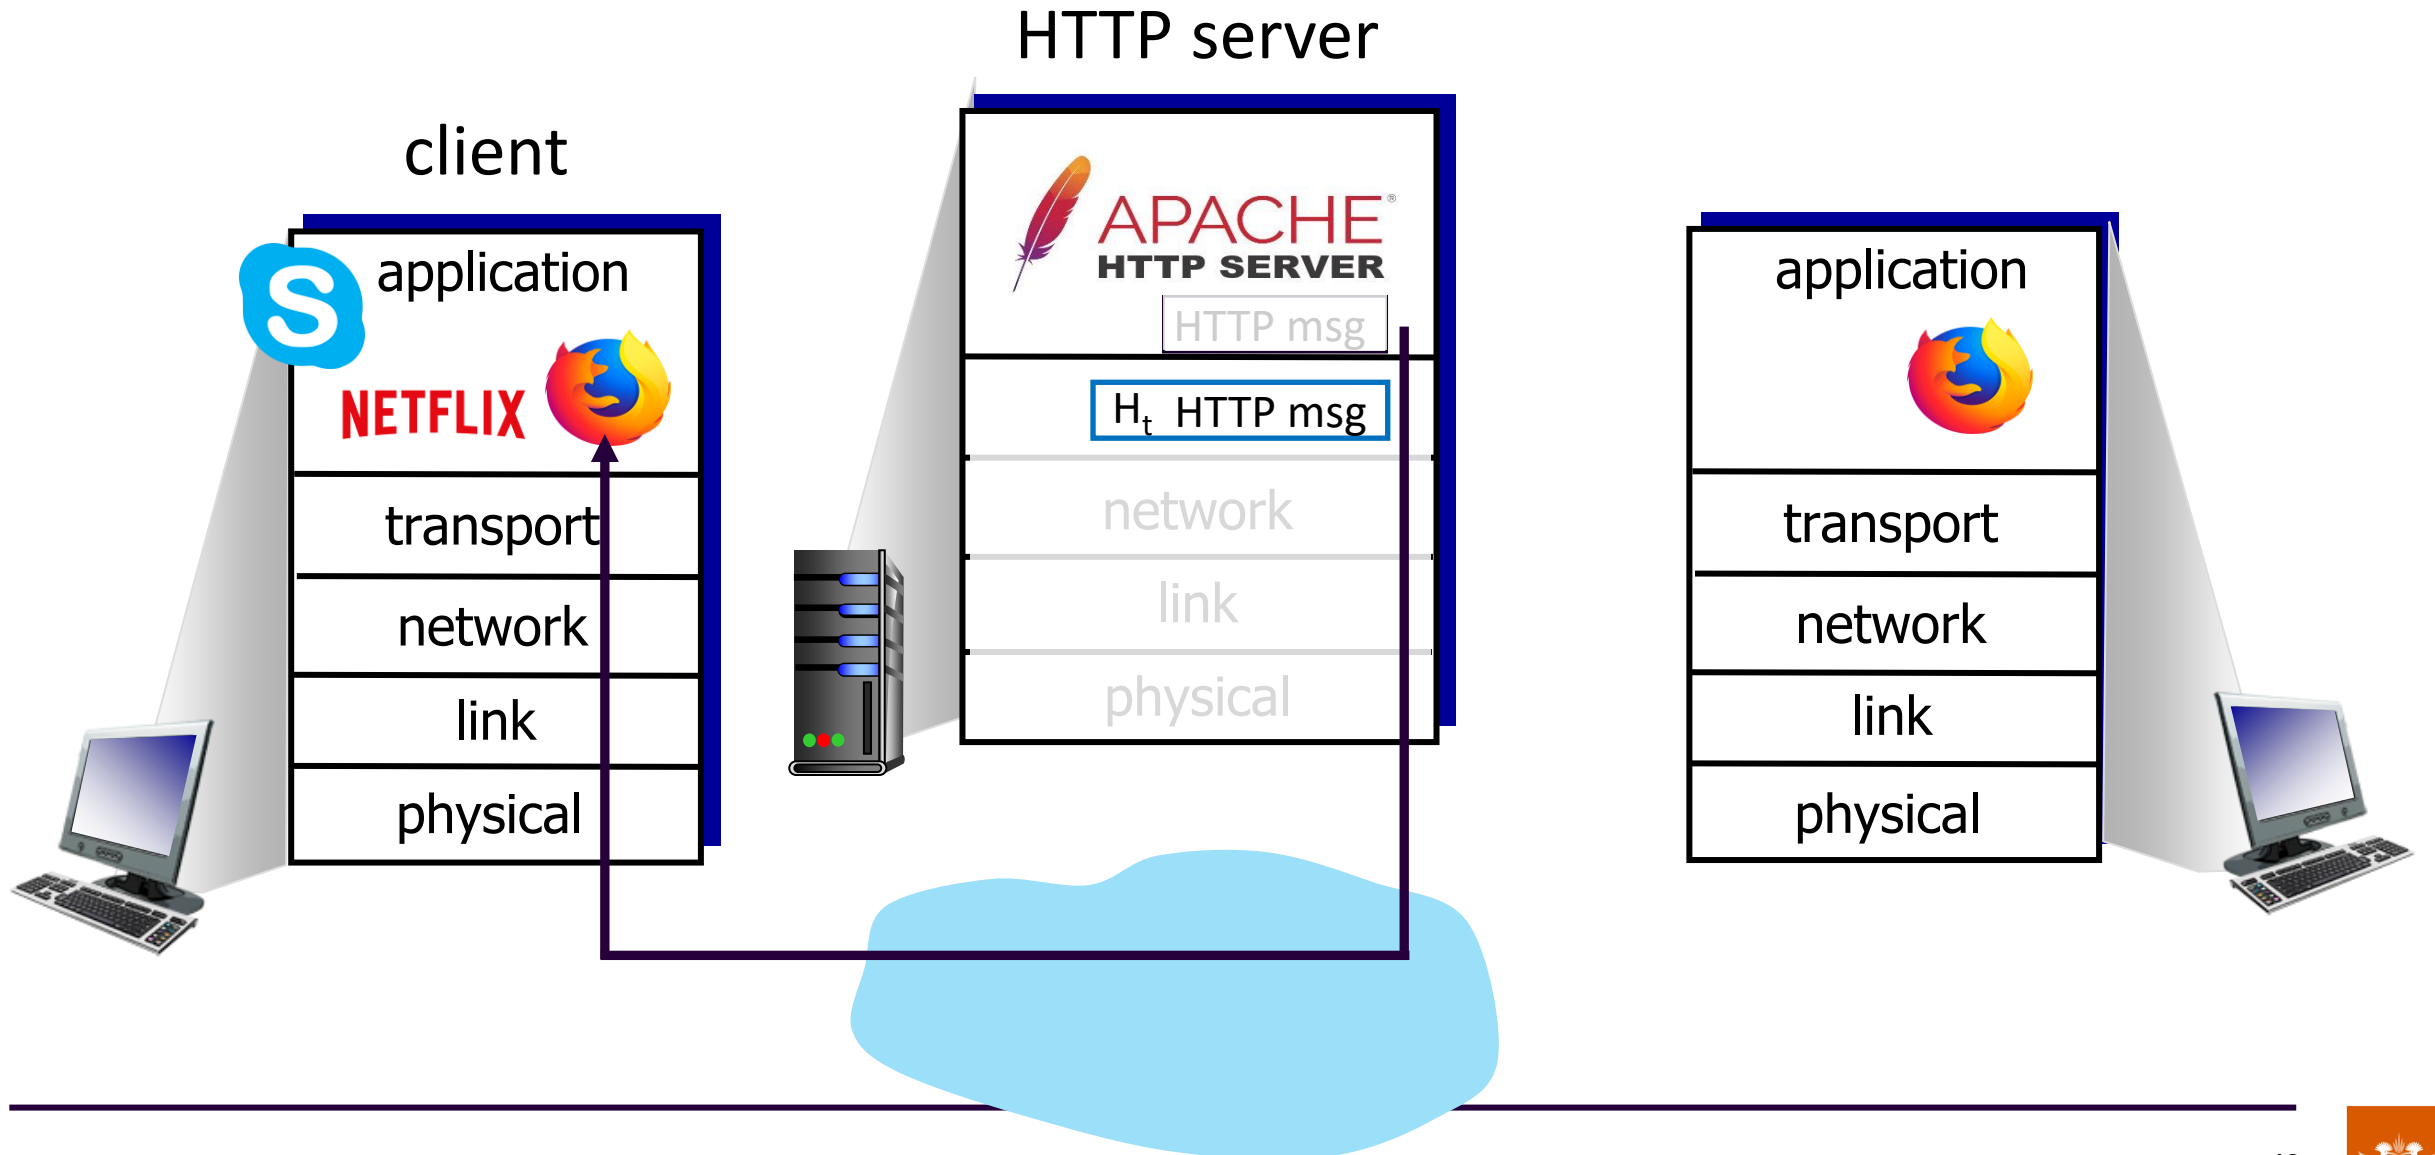
- segment structure
- reliable data transfer
- flow control
- connection management

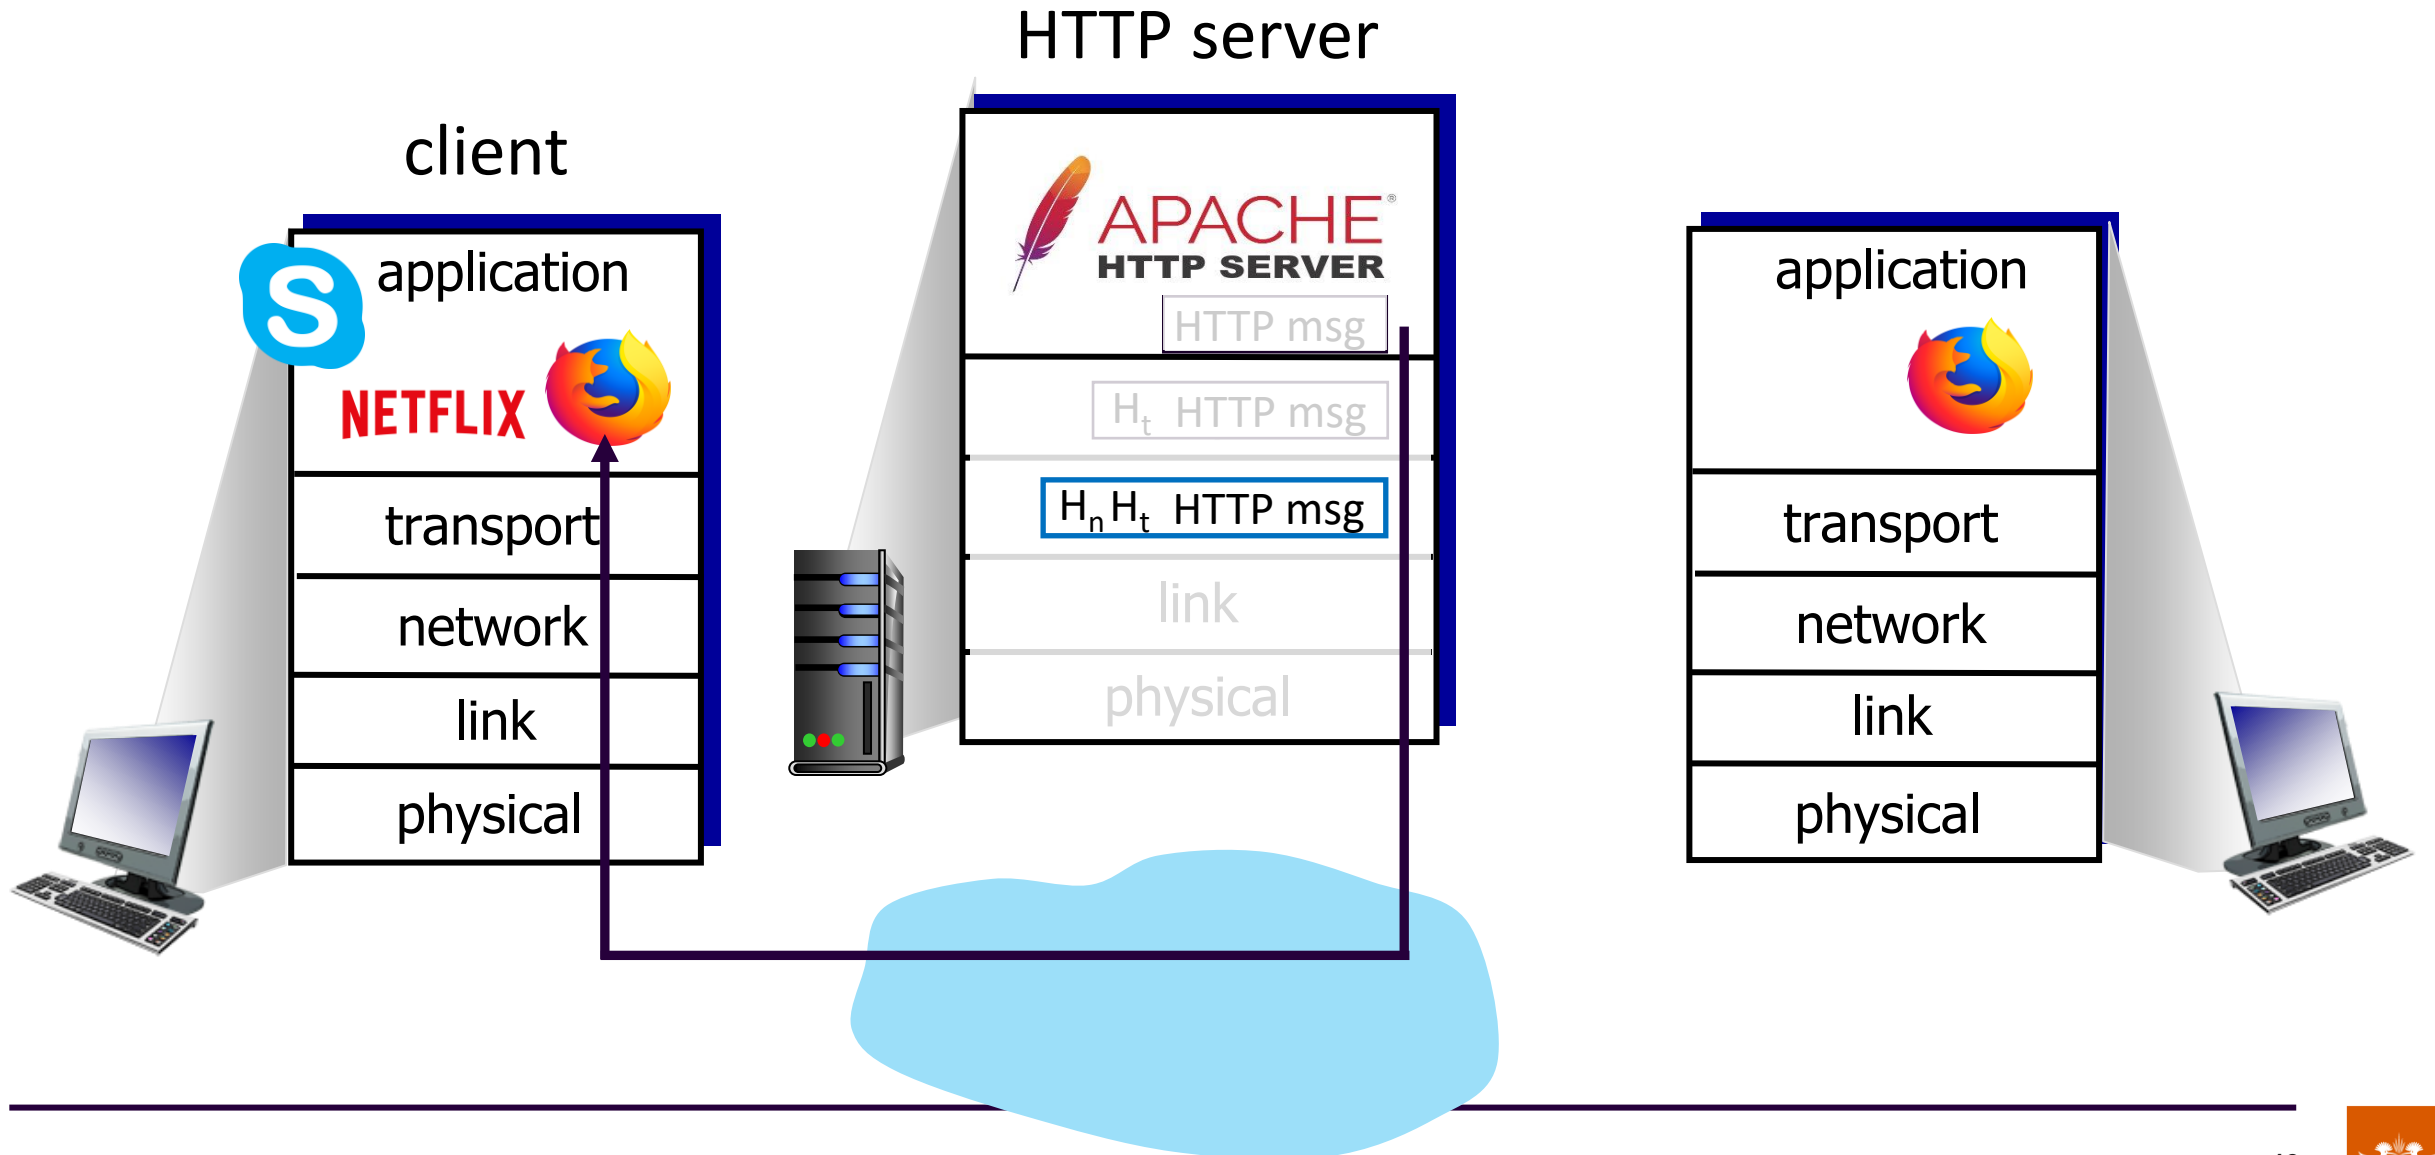
3.6 Principles of congestion control

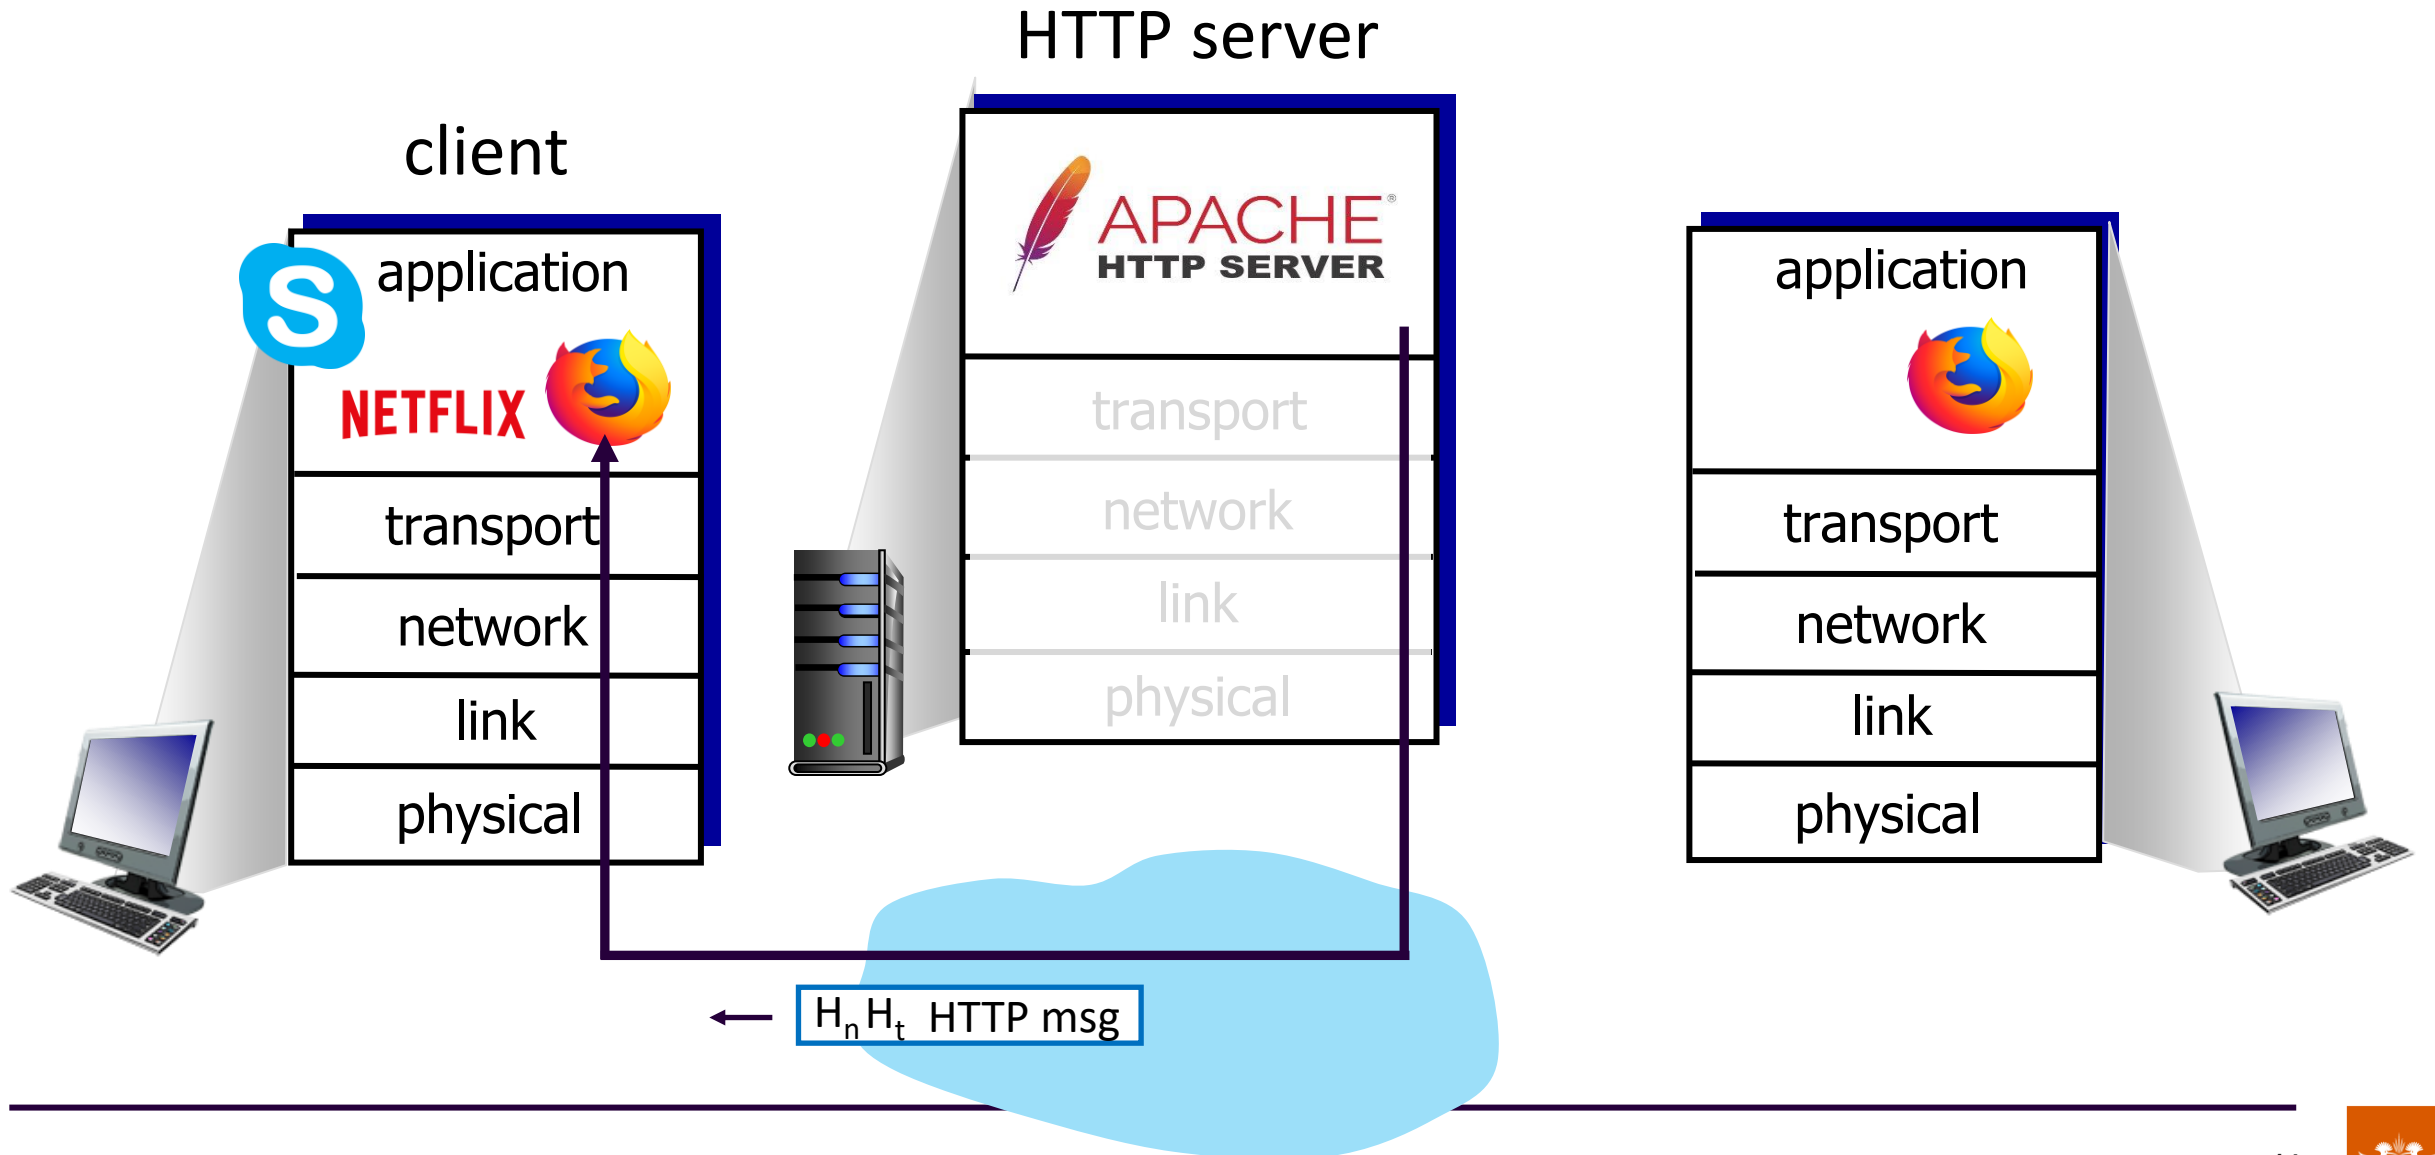
3.7 TCP congestion control

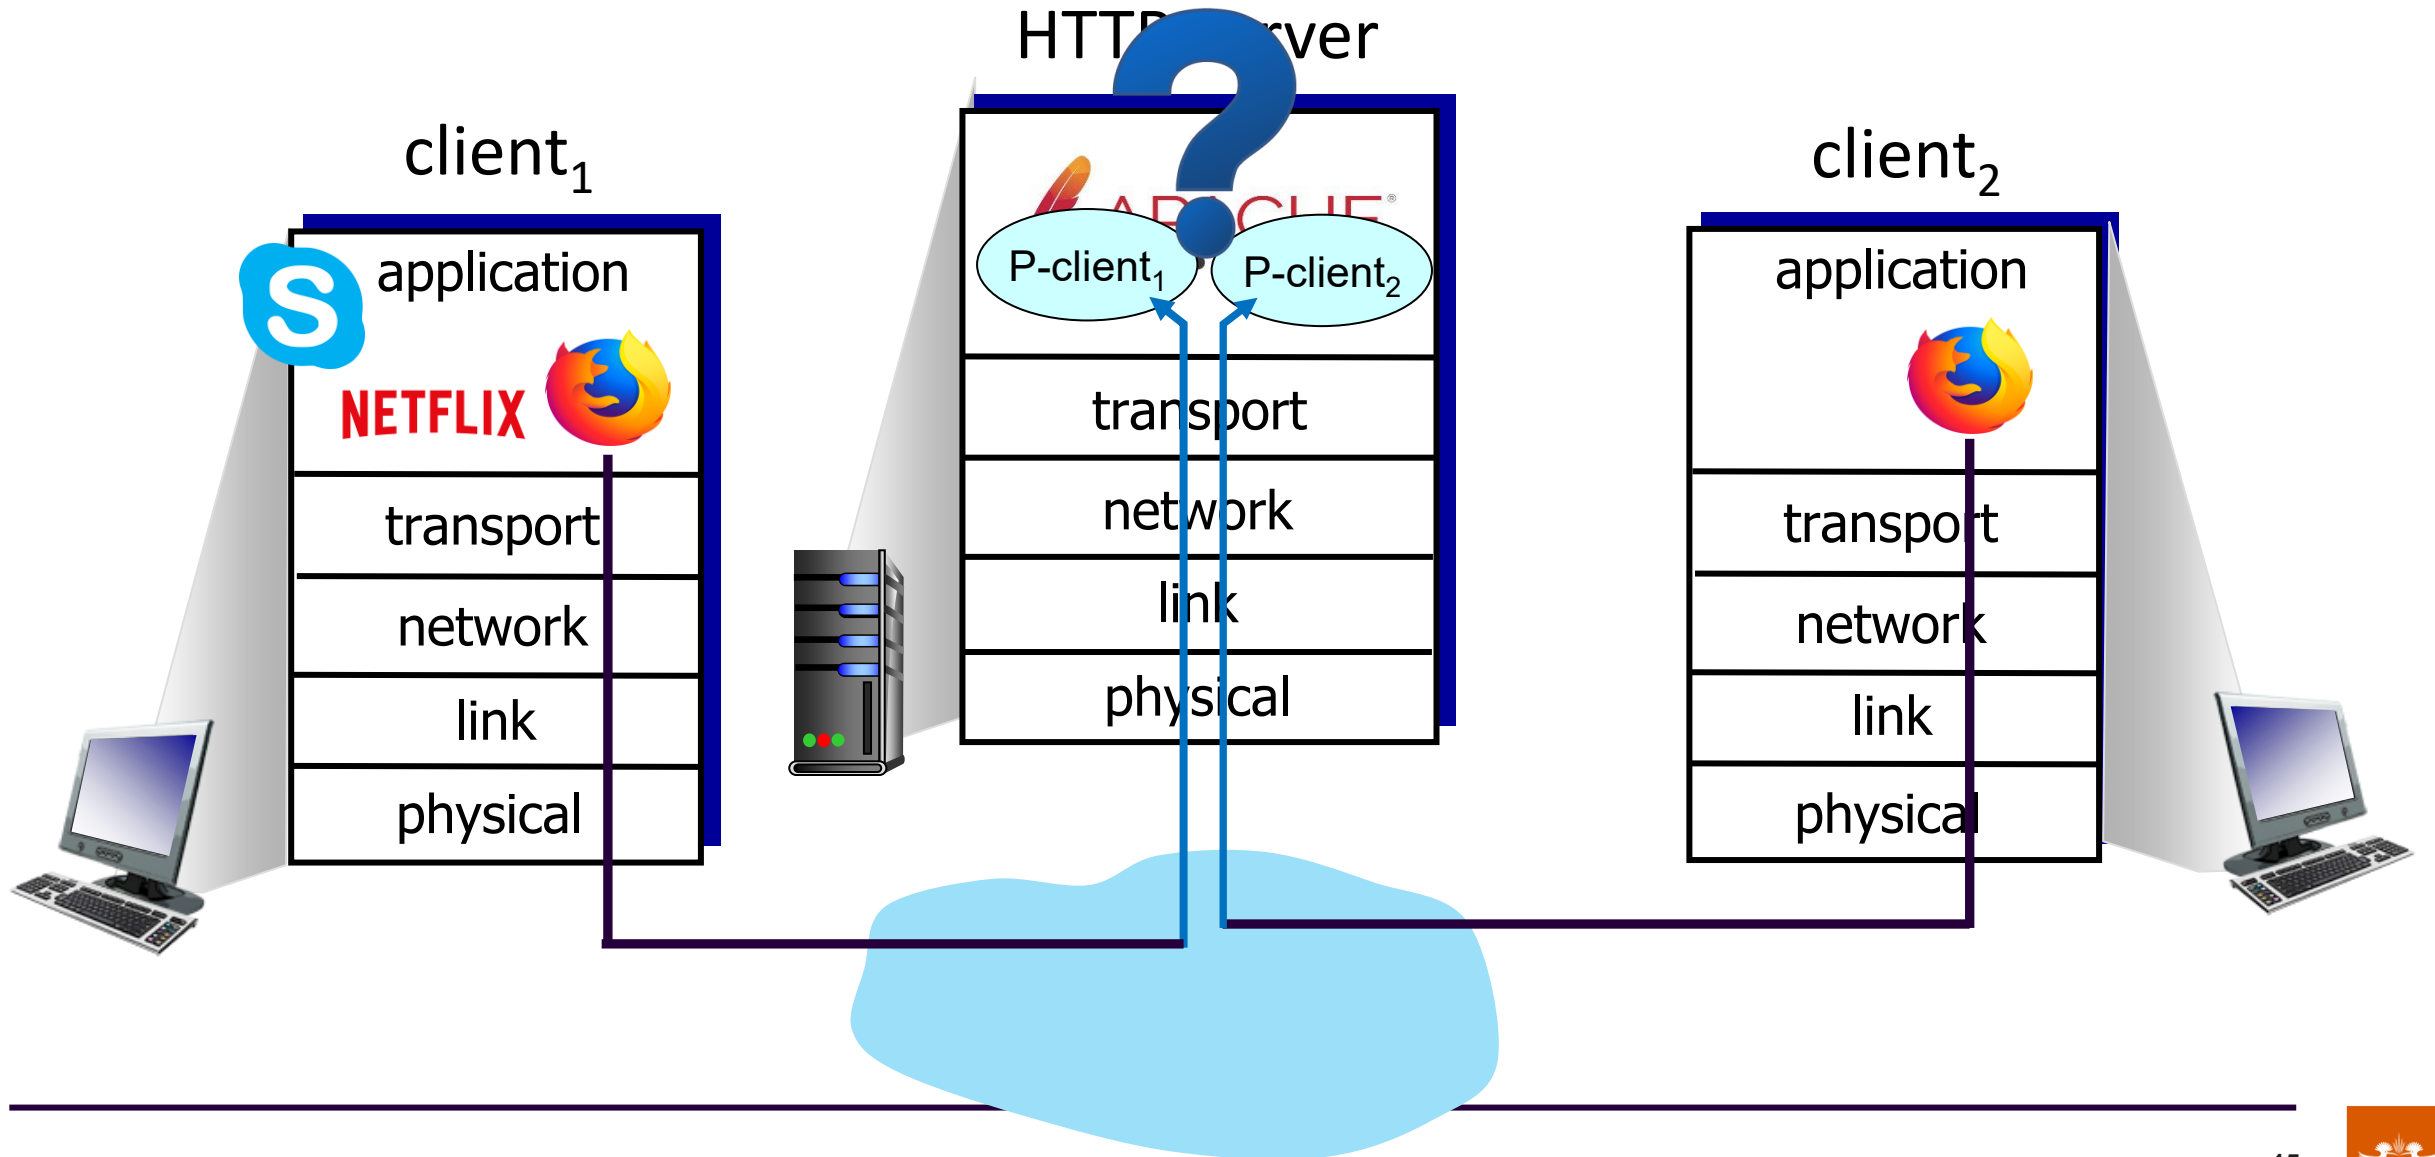












Multiplexing/demultiplexing

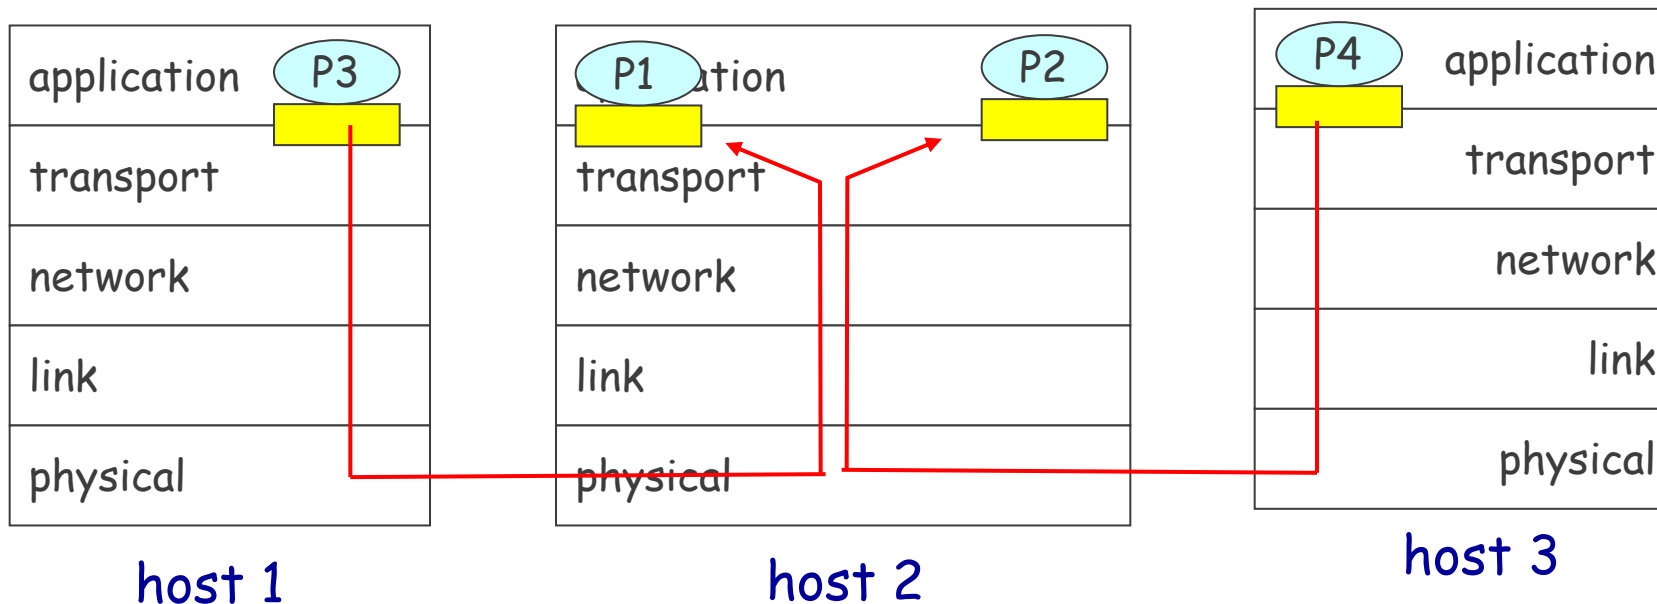
Demultiplexing at rcv host:

delivering received segments
to correct socket

Multiplexing at send host:

gathering data from multiple
sockets, enveloping data with
header (later used for
demultiplexing)

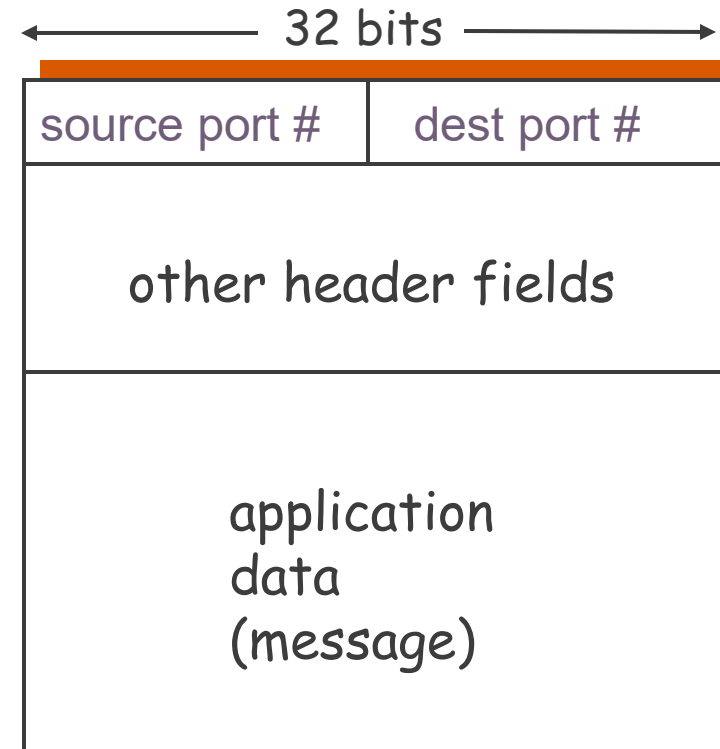
 = socket  = process



How demultiplexing works

- host receives IP datagrams
 - each datagram has source IP address, destination IP address
 - each datagram carries 1 transport-layer segment
 - each segment has source, destination port number
- host uses IP addresses & port numbers to direct segment to appropriate socket

TCP/UDP segment format



Connectionless demultiplexing

Recall:

- when creating socket, must specify *host-local* port #:

```
DatagramSocket mySocket1  
= new DatagramSocket(12534);
```

- when creating datagram to send into UDP socket, must specify
 - destination IP address
 - destination port #

when receiving host receives *UDP* segment:

- checks destination port # in segment
- directs UDP segment to socket with that port #

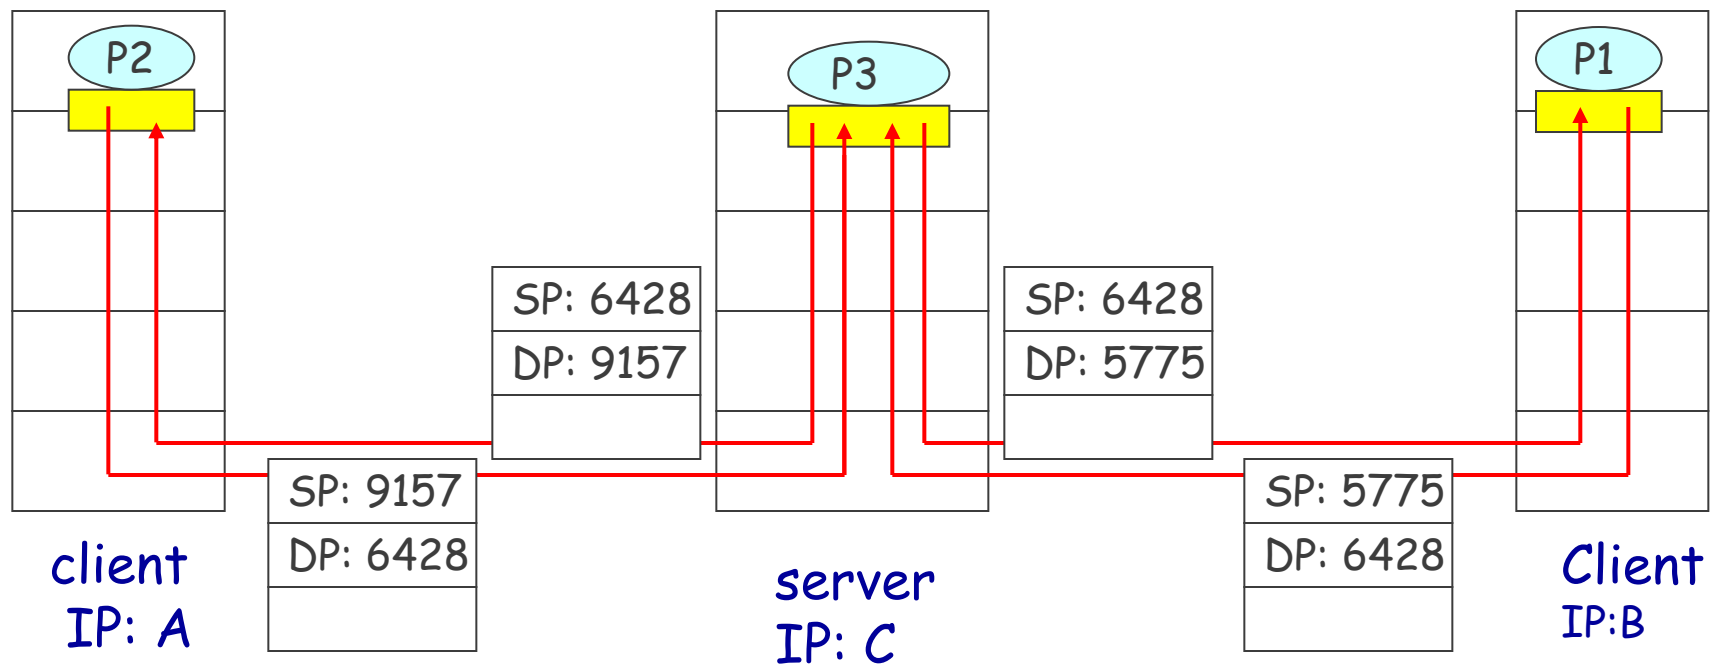


IP/UDP datagrams with *same dest. port #*, but different source IP addresses and/or source port numbers will be directed to *same socket* at receiving host



Connectionless demux (cont)

```
DatagramSocket serverSocket = new DatagramSocket(6428);
```



SP provides “return address”

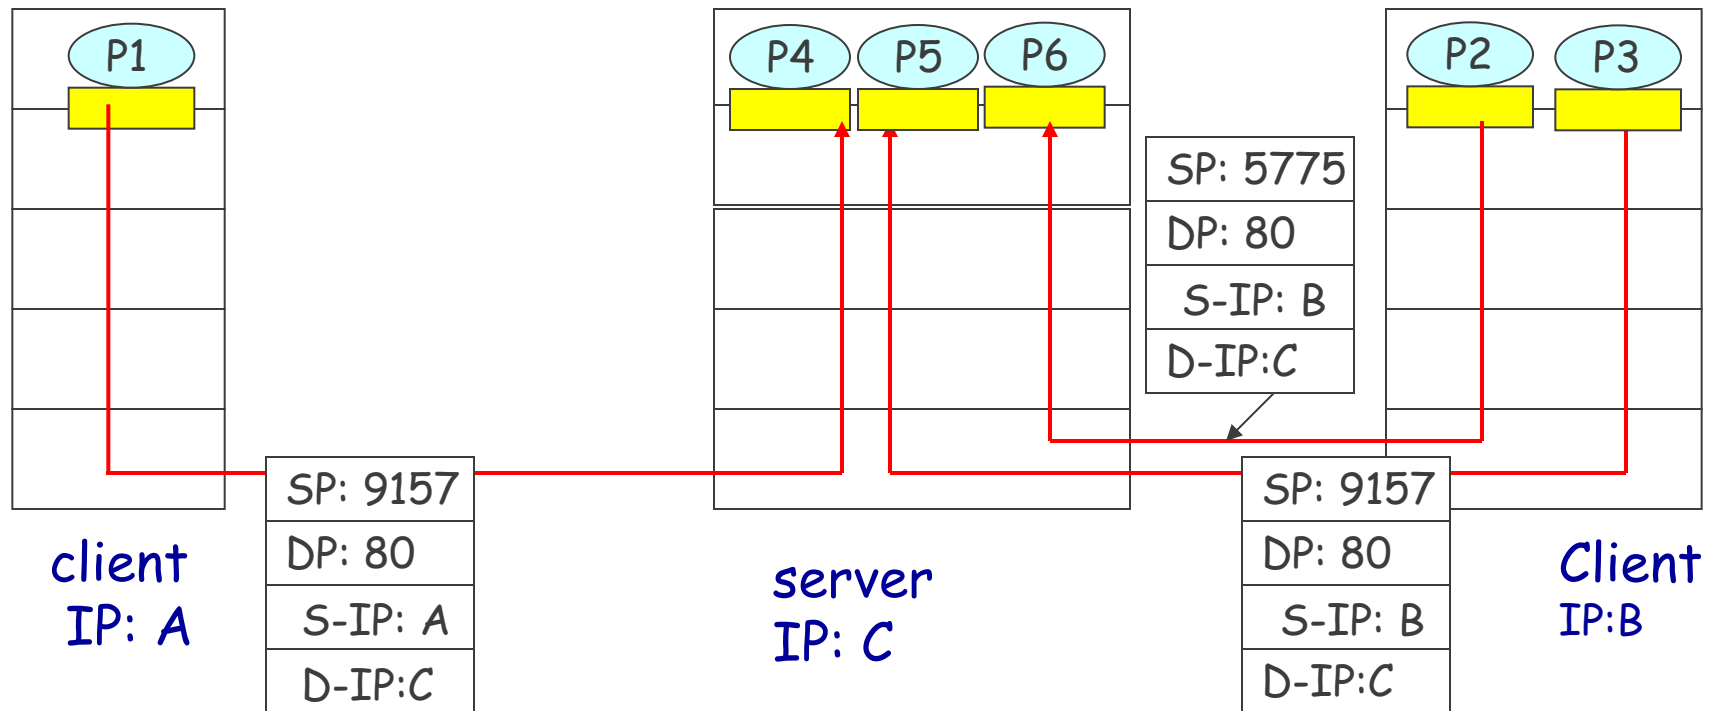


Connection-oriented demux

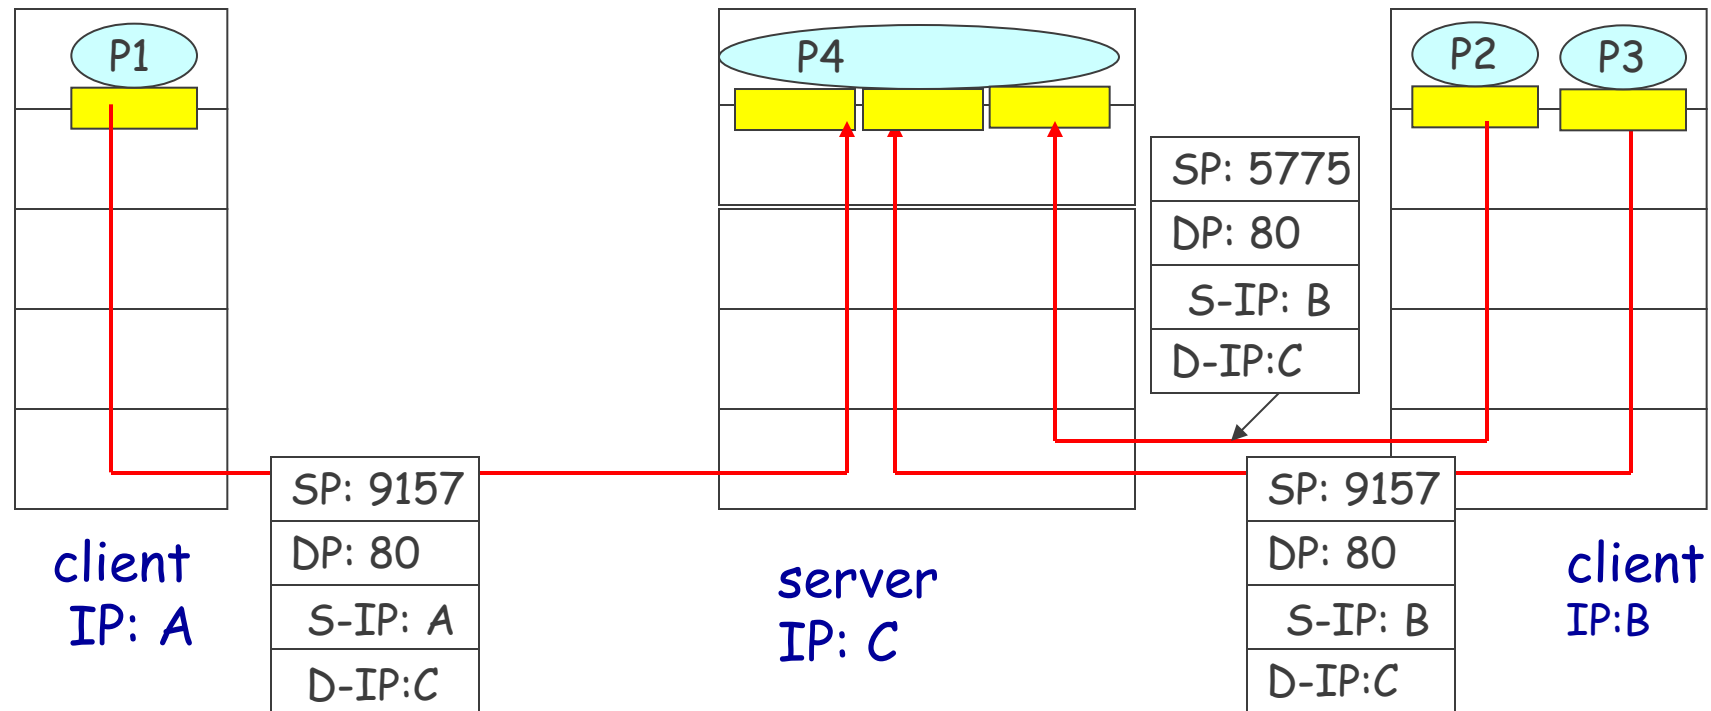
- TCP socket identified by 4-tuple:
 - **source IP address**
 - **source port number**
 - **dest IP address**
 - **dest port number**
- recv host uses all four values to direct segment to appropriate socket
- server host may support many simultaneous TCP sockets:
 - each socket identified by its own 4-tuple
- web servers have different sockets for each connecting client
 - non-persistent HTTP will have different socket for each request



Connection-oriented demux (cont)



Connection-oriented demux: Threaded Web Server



Summary

- Multiplexing, demultiplexing: based on segment, datagram header field values
- **UDP:** demultiplexing using destination port number (only)
- **TCP:** demultiplexing using 4-tuple: source and destination IP addresses, and port numbers
- Multiplexing/demultiplexing happen at *all* layers



Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- **Connectionless transport: UDP**
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer functionality



UDP: User Datagram Protocol [RFC 768]

- “no frills,” “bare bones” Internet transport protocol
- “best effort” service, UDP segments may be:
 - lost
 - delivered out of order to app
- **connectionless:**
 - no handshaking between UDP sender, receiver
 - each UDP segment handled independently of others

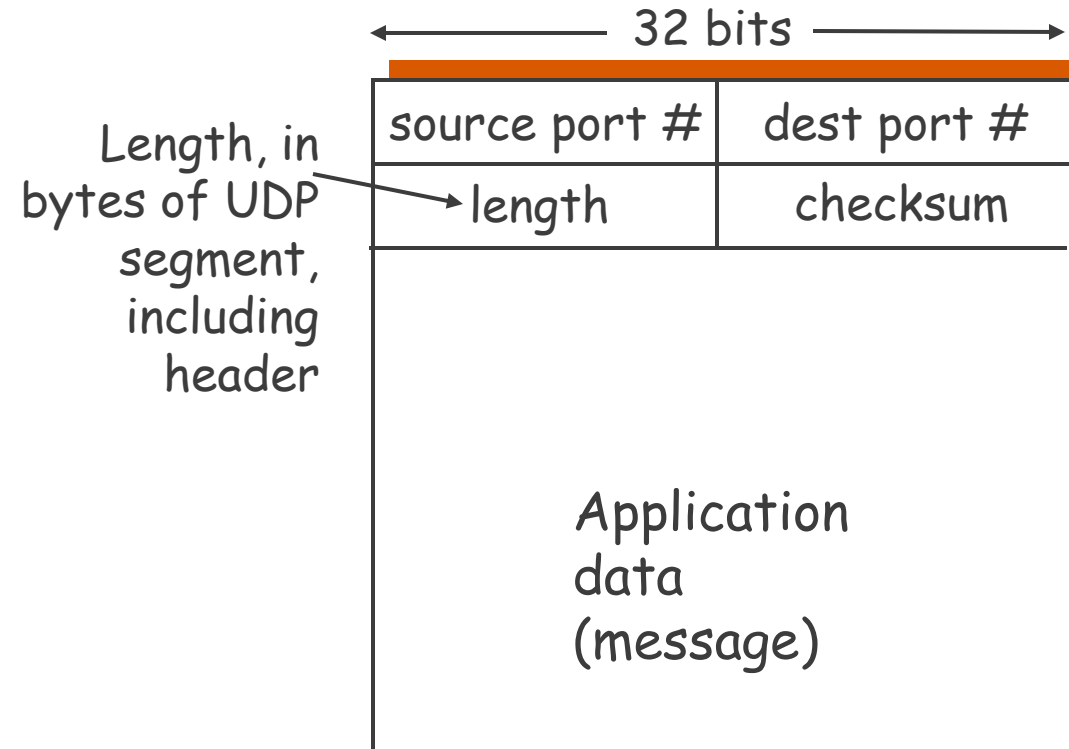
Why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small segment header
- no congestion control: UDP can blast away as fast as desired



UDP: more

- often used for streaming multimedia apps
 - loss tolerant
 - rate sensitive
- other UDP uses
 - DNS
 - SNMP
- reliable transfer over UDP: add reliability at application layer
 - application-specific error recovery!



UDP segment format



UDP checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment

Sender:

- treat segment contents as sequence of 16-bit integers
- checksum: addition (1's complement sum) of segment contents
- sender puts checksum value into UDP checksum field

Receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - NO - error detected
 - YES - no error detected. *But maybe errors nonetheless? More later*



Internet Checksum Example

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result

Example: add two 16-bit integers

		1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
		1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
		<hr/>															
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1	1
		<hr/>															
sum		1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum		0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

